



# Building Knowledge Graphs

## Building Knowledge Graphs mind map





## Flashcards easy

| No | Question                                                                                 | Answer                                                                             |
|----|------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| 1  | What is the primary characteristic that transforms a plain graph into a knowledge graph? | The application of an organizing principle that provides contextual understanding. |
| 2  | In the property graph model, what do nodes represent?                                    | Entities in the domain.                                                            |

| No | Question                                                                                        | Answer                                                                                                                |
|----|-------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
| 3  | What elements in a property graph represent how entities interrelate?                           | Relationships (or edges).                                                                                             |
| 4  | What are the key-value pairs stored on nodes or relationships called?                           | Properties.                                                                                                           |
| 5  | True or False: In a property graph, relationships can exist without an end node.                | False; relationships never dangle and must always have a start and an end node.                                       |
| 6  | Which graph model element is used to declare a node's purpose, such as 'Customer' or 'Product'? | Labels.                                                                                                               |
| 7  | Who invented graph theory in the 18th century to solve the Seven Bridges of Königsberg problem? | Leonhard Euler.                                                                                                       |
| 8  | What is the difference between a graph and a chart?                                             | Graphs show how things connect (networks), while charts visualize data distributions or functions (e.g., histograms). |

| No | Question                                                                               | Answer                                                                                                           |
|----|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| 9  | Why is the property graph model described as 'schema-ish'?                             | It is flexible and doesn't enforce a rigid schema but allows optional constraints on labels and properties.      |
| 10 | What is the function of an 'organizing principle' in a knowledge graph?                | It provides a layer of structure and context that allows humans and computers to reason about the data.          |
| 11 | How does a taxonomy differ from a plain graph?                                         | A taxonomy organizes categories into a broader-narrower or generalization-specialization hierarchy.              |
| 12 | Which hierarchical relationship is typically used in a taxonomy to connect categories? | SUBCATEGORY_OF (or NARROWER_THAN/SUBCLASS_OF).                                                                   |
| 13 | What is an ontology in the context of knowledge graphs?                                | A classification scheme that defines complex categories and varied relationship types beyond simple hierarchies. |
| 14 | Give an example of a horizontal (cross-cutting) relationship defined in an ontology.   | PART_OF, COMPATIBLE_WITH, or DEPENDS_ON.                                                                         |

| No | Question                                                                                 | Answer                                                                                                           |
|----|------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| 15 | What is the 'just-enough semantics' principle?                                           | The practice of introducing only the semantic metadata required for current use cases to avoid over-engineering. |
| 16 | Name one standard language used to formally define ontologies.                           | Web Ontology Language (OWL) or RDF Schema.                                                                       |
| 17 | What does SKOS stand for in the context of taxonomical classification?                   | Simple Knowledge Organization System.                                                                            |
| 18 | Term: Cypher                                                                             | Definition: A declarative, visual pattern-matching query language for graph databases.                           |
| 19 | In Cypher syntax, what symbols represent a node?                                         | Parentheses ( ).                                                                                                 |
| 20 | In Cypher syntax, what symbols represent a relationship?                                 | Brackets within an arrow, such as -[ ]->.                                                                        |
| 21 | Which Cypher keyword is used to insert data into a graph, even if it creates duplicates? | CREATE                                                                                                           |

| No | Question                                                                                            | Answer                                                                                                        |
|----|-----------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| 22 | Which Cypher keyword acts as a mix of MATCH and CREATE, ensuring no duplicate patterns are created? | MERGE                                                                                                         |
| 23 | What is the purpose of the DETACH DELETE command in Cypher?                                         | It removes a node and all of its associated relationships simultaneously.                                     |
| 24 | In graph database performance, query latency is typically proportional to what factor?              | The amount of the graph traversed (not the total size of the database).                                       |
| 25 | What is a 'graph local' query?                                                                      | A query anchored to a specific, known node in the graph.                                                      |
| 26 | What is a 'graph global' query?                                                                     | A query that processes the entire graph or large parts of it without being bound to a specific starting node. |
| 27 | Which Cypher keyword is used to visualize a query plan without executing it?                        | EXPLAIN                                                                                                       |

| No | Question                                                                                                  | Answer                                                                                                                                 |
|----|-----------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| 28 | Which Cypher keyword instruments and runs a query to provide runtime performance feedback?                | PROFILE                                                                                                                                |
| 29 | Term: Index-free adjacency                                                                                | Definition: A storage architecture where nodes and relationships point directly to each other via offsets, enabling $O(1)$ traversals. |
| 30 | How is the time complexity of accessing a node or relationship in Neo4j's traversal engine described?     | Constant time, or $O(1)$ .                                                                                                             |
| 31 | In Neo4j, property storage is described as linear time, or $O(N)$ .                                       | $O(N)$                                                                                                                                 |
| 32 | What is the purpose of a 'write-ahead log' in graph databases?                                            | To ensure ACID transactions and provide fault tolerance by persisting updates to disk before applying them.                            |
| 33 | Which tool provides a visual interface to map CSV files onto a graph domain model for loading into Neo4j? | Neo4j Data Importer                                                                                                                    |



| No | Question                                                                                                                     | Answer                                                                                                     |
|----|------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| 34 | Which Cypher clause transforms a list of values into individual rows for subsequent processing?                              | UNWIND                                                                                                     |
| 35 | What is the primary command for importing data from web-based or local CSV files while the database is online?               | LOAD CSV                                                                                                   |
| 36 | For extremely large datasets (billions of records), which Neo4j command-line tool provides high-performance offline loading? | neo4j-admin import                                                                                         |
| 37 | In the context of information systems, what is a 'data fabric'?                                                              | A general-purpose data access layer that offers a unified, connected view of data across underlying silos. |

| No | Question                                                                                                     | Answer                                                                           |
|----|--------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|
| 38 | Which component of a data fabric often uses a knowledge graph to curate data across different storage types? | The enterprise index (or master data layer).                                     |
| 39 | In Neo4j driver development, should you create one Driver object per query or one per application lifetime?  | One per application lifetime (Driver objects are expensive).                     |
| 40 | Which Neo4j driver object is considered 'cheap' and intended for individual units of work?                   | Session                                                                          |
| 41 | What is the purpose of 'parameterizing' queries in a database driver?                                        | To increase performance by reusing query plans and to prevent injection attacks. |
| 42 | Which Neo4j feature allows multiple physical databases to be treated as a single unified virtual graph?      | Composite Databases (Graph Federation).                                          |

| No | Question                                                                                                               | Answer                                                                                                             |
|----|------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| 43 | What is the APOC library in the Neo4j ecosystem?                                                                       | A library containing hundreds of procedures and functions for common tasks like data loading and virtualization.   |
| 44 | What is 'data virtualization' in a knowledge graph context?                                                            | Retrieving data from external sources on demand and presenting it as virtual nodes/relationships within the graph. |
| 45 | Which API framework is commonly used to expose knowledge graphs to frontend applications via a schema-based interface? | GraphQL                                                                                                            |
| 46 | How can a knowledge graph integrate with real-time event streams like Apache Kafka?                                    | Using a Kafka Connect plug-in to act as a source or sink for data events.                                          |
| 47 | Which distributed data processing framework uses a specialized connector to treat Neo4j as a source or sink?           | Apache Spark                                                                                                       |
| 48 | What is the goal of graph data science?                                                                                | To gain insight into a knowledge graph's structure using graph algorithms.                                         |

| No | Question                                                                               | Answer                                                                                             |
|----|----------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|
| 49 | Name the three broad categories of graph algorithms described in the text.             | Statistical, Analytical, and Machine Learning.                                                     |
| 50 | What do 'Centrality' algorithms measure in a graph?                                    | The influence or importance of individual nodes (e.g., bridges or bottlenecks).                    |
| 51 | Which algorithm class identifies tightly knit groups of nodes within a larger network? | Community Detection                                                                                |
| 52 | What are the four phases of executing a graph algorithm in Neo4j Graph Data Science?   | Read projection, Load projection, Execute algorithm, Store results.                                |
| 53 | What is a 'graph projection' in Neo4j Graph Data Science?                              | An in-memory, compressed representation of a specific subgraph optimized for parallel computation. |
| 54 | Define 'Betweenness Centrality'.                                                       | A metric representing the number of shortest paths that pass through a given node.                 |

| No | Question                                                                                                           | Answer                                                                                                        |
|----|--------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| 55 | In a production environment, why are 'secondary' servers often used for graph data science?                        | To isolate heavy computational workloads from the transactional processing of 'primary' servers.              |
| 56 | What mechanism ensures a user sees their own writes even when querying an asynchronously updated secondary server? | Causal barrier (using Bookmarks).                                                                             |
| 57 | In graph machine learning, what is 'in-graph machine learning'?                                                    | Using a knowledge graph to compute and predict its own missing elements, like relationships or labels.        |
| 58 | What is 'graph feature engineering'?                                                                               | Extracting topological features (like centrality or embeddings) from a graph to train external ML models.     |
| 59 | What is a 'link prediction' algorithm used for?                                                                    | Determining the likelihood that a relationship should exist between two nodes based on the existing topology. |
| 60 | Term: Node Embedding                                                                                               | Definition: An algorithmic encoding of a node's local topology into a fixed-length numerical vector.          |

| No | Question                                                                       | Answer                                                                                           |
|----|--------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------|
| 61 | Which Neo4j algorithm is frequently used for high-performance node embeddings? | FastRP (Fast Random Projection).                                                                 |
| 62 | What are the components of a link prediction pipeline in Neo4j GDS?            | Node properties, link features, split configuration, model candidates, and training.             |
| 63 | What does a 'metadata knowledge graph' map?                                    | The location of data, the systems that process it (pipelines), and the final consumers (sinks).  |
| 64 | In a metadata graph, what is a 'Data Sink'?                                    | A terminal consumer of data, such as a dashboard visualization or an ML training set.            |
| 65 | What is 'data lineage' in a metadata graph context?                            | Tracing the flow of data back to its original source platform across various tasks and datasets. |
| 66 | What is 'entity resolution'?                                                   | The process of determining if two separate records represent the same real-world thing.          |
| 67 | What are the three fundamental activities in graph-based entity resolution?    | Data preparation, Entity matching, and Master entity curation.                                   |

| No | Question                                                                                                      | Answer                                                                                                                    |
|----|---------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| 68 | In entity resolution, what is a 'blocking key'?                                                               | A value used to partition data into groups so that only likely matches are compared, reducing the search space.           |
| 69 | What relationship type is used to connect nodes that definitively represent the same entity after resolution? | SAME_AS                                                                                                                   |
| 70 | What is the purpose of a 'SIMILAR' relationship in entity resolution?                                         | To connect nodes that meet a probabilistic threshold for being the same entity, often annotated with a similarity score.  |
| 71 | Which string similarity metric is mentioned in the text for matching weak identifiers like names?             | Jaro–Winkler distance.                                                                                                    |
| 72 | In the entity resolution process, what is a 'strong identifier'?                                              | A unique, unambiguous value like a Social Security number or passport number.                                             |
| 73 | Why is 'index-free adjacency' more efficient for graph traversals than using global join indexes?             | It uses local pointers to move between records, avoiding the $O(\log N)$ cost of searching a global index for every step. |

| No | Question                                                                                                            | Answer                                                                                                         |
|----|---------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| 74 | In a Neo4j cluster, which server type is responsible for handling transactional writes?                             | Primary server.                                                                                                |
| 75 | Which GDS procedure mode writes results directly back to the nodes and relationships in the database?               | write                                                                                                          |
| 76 | Which GDS procedure mode streams results back to the caller without changing the graph stored in memory or on disk? | stream                                                                                                         |
| 77 | Which GDS procedure mode enriches only the in-memory graph projection?                                              | mutate                                                                                                         |
| 78 | What is the primary benefit of treating metadata and data as separate layers within the same knowledge graph?       | It allows users to search for data based on business concepts and provides built-in provenance and governance. |



## Flashcards medium

| No | Question                                                                             | Answer                                                                                                                |
|----|--------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
| 1  | What was Leonhard Euler's primary insight regarding the Königsberg bridge problem?   | The physical geography was irrelevant; only the logical connections between landmasses mattered.                      |
| 2  | In the context of knowledge graphs, how do 'charts' differ from 'graphs'?            | Charts visualize data trends (like histograms), while graphs model networks of entities and relationships.            |
| 3  | What is the definition of a 'knowledge graph'?                                       | An interlinked set of facts describing real-world entities and their relations in a machine-understandable format.    |
| 4  | Why is an 'organizing principle' essential for a knowledge graph?                    | It provides a layer of structure and rules that allow humans or computer systems to reason about the underlying data. |
| 5  | In a Labeled Property Graph (LPG), what do 'nodes' represent?                        | Entities within a specific domain.                                                                                    |
| 6  | What are the three core elements that can be stored on a 'node' in a property graph? | Labels, properties (key-value pairs), and relationships.                                                              |
| 7  | What is the function of a 'label' in a property graph?                               | It declares the node's purpose or role in the graph, such as 'Customer' or 'Product'.                                 |
| 8  | What rule governs 'relationships' in a property graph to ensure data integrity?      | They never 'dangle'; they must always have a start node and an end node.                                              |

| No | Question                                                         | Answer                                                                                                                              |
|----|------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| 9  | In the LPG model, how are 'properties' defined?                  | They are key-value pairs representing entity data, such as a name or date of birth.                                                 |
| 10 | The property graph model is described as 'schema-ish' because _. | Constraints are applied only where needed rather than locking down the entire data model.                                           |
| 11 | What distinguishes a 'plain old graph' from a 'knowledge graph'? | Plain old graphs encode their interpretation in application logic, while knowledge graphs include an explicit organizing principle. |
| 12 | Term: Taxonomy                                                   | Definition: A classification scheme that organizes categories in a broader-narrower or generalization-specialization hierarchy.     |
| 13 | How does an 'ontology' differ from a 'taxonomy'?                 | Ontologies define complex, non-hierarchical relationships and semantics, such as 'PART_OF' or 'COMPATIBLE_WITH'.                    |
| 14 | What is the 'just-enough semantics' principle?                   | The practice of introducing semantic metadata only as current use cases demand to avoid overengineering.                            |
| 15 | Term: Semantic Bridge                                            | Definition: An ontology that defines cross-equivalence to link the same concepts across different systems.                          |
| 16 | What is the primary design goal of the Cypher query language?    | To be a visual, humane, and expressive declarative pattern-matching language.                                                       |

| No | Question                                                                                        | Answer                                                                                                                              |
|----|-------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| 17 | In Cypher syntax, what do parentheses <code>()</code> represent?                                | Nodes.                                                                                                                              |
| 18 | What is the purpose of the <code>MATCH</code> keyword in Cypher?                                | To compare user-defined patterns against the underlying knowledge graph data.                                                       |
| 19 | How does the <code>CREATE</code> keyword behave in Cypher compared to <code>MERGE</code> ?      | <code>CREATE</code> always inserts new data, while <code>MERGE</code> matches existing patterns or creates them if they are absent. |
| 20 | The Cypher keyword used to remove a node and all its attached relationships is <code>_</code> . | <code>DETACH DELETE</code>                                                                                                          |
| 21 | What does a 'graph local query' mean?                                                           | A query anchored to a specific, known node in the graph rather than searching the entire database.                                  |
| 22 | How does graph database size affect query latency in systems like Neo4j?                        | Latency is proportional to the amount of the graph traversed, not the total size of the database.                                   |
| 23 | Term: Index-Free Adjacency                                                                      | Definition: A storage technique where nodes store direct pointers to adjacent records, enabling $O(1)$ traversals.                  |
| 24 | What is the time complexity of reading a single property in Neo4j?                              | It is $O(N)$ , where $N$ is the number of properties on that specific node or relationship.                                         |
| 25 | Why is 'scale-up' (adding RAM/CPU) generally preferred over 'scale-out' for graph processing?   | Scale-up maintains data locality in memory, avoiding the high latency of network-based distributed traversals.                      |

| No | Question                                                                                                               | Answer                                                                                                                         |
|----|------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| 26 | What role does a 'write-ahead log' play in Neo4j internals?                                                            | It ensures durability by persisting intended updates to disk before they are applied to the graph.                             |
| 27 | What is the main advantage of the 'Neo4j Data Importer'?                                                               | It provides a no-code visual interface to map CSV files to a graph domain model.                                               |
| 28 | When using <code>LOAD CSV</code> in Cypher, what does the <code>WITH HEADERS</code> clause do?                         | It allows the query to refer to CSV columns by their names rather than numerical indices.                                      |
| 29 | Which Neo4j tool provides the fastest possible method for initial bulk data loading?                                   | neo4j-admin import                                                                                                             |
| 30 | Term: Data Fabric                                                                                                      | Definition: A connected view of data across an organization that abstracts away the details of underlying third-party systems. |
| 31 | Regarding Neo4j drivers, why should a 'Driver' object be instantiated only once for an application's lifetime?         | Driver objects are expensive to create because they manage authentication and network connection pools.                        |
| 32 | In a Neo4j cluster, what is the purpose of distinguishing between 'read' and 'write' transactions in application code? | It allows the driver to make better routing decisions to maximize cluster throughput.                                          |
| 33 | What is 'Graph Federation'?                                                                                            | A technique using composite databases to provide a single unified entry point for multiple independent graph sources.          |

| No | Question                                                            | Answer                                                                                                                          |
|----|---------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| 34 | How does 'Data Virtualization' with APOC benefit a knowledge graph? | It allows the graph to retrieve and present external data (like SQL or JSON) as virtual nodes without copying the data.         |
| 35 | Term: Graph Data Science (GDS)                                      | Definition: The application of graph algorithms to uncover structural insights, patterns, or features within a knowledge graph. |
| 36 | What are 'Statistical' graph algorithms primarily used for?         | Providing high-level metrics like node counts, relationship types, and degree distributions.                                    |
| 37 | What is the goal of 'Community Detection' algorithms?               | To partition the graph into groups of nodes that are more densely connected to each other than to the rest of the network.      |
| 38 | In GDS, what does 'Centrality' (or Influence) measure?              | The relative importance or 'bridging' role of a node within the network topology.                                               |
| 39 | What insight does the 'Betweenness Centrality' algorithm provide?   | It identifies nodes that act as 'bridges' or bottlenecks by counting the shortest paths passing through them.                   |
| 40 | What is the first step in the Neo4j GDS workflow?                   | Creating a graph projection (loading a subgraph into memory for processing).                                                    |
| 41 | When executing a GDS algorithm, what does the 'mutate' mode do?     | It writes the results of the algorithm back into the in-memory projection rather than the persistent database.                  |

| No | Question                                                                     | Answer                                                                                                                   |
|----|------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| 42 | How does 'In-graph Machine Learning' differ from traditional ML?             | It allows a model to predict its own evolution (like missing links) by computing directly over its topology.             |
| 43 | Term: Graph Feature Engineering                                              | Definition: The process of extracting topological metrics (like PageRank or embeddings) from a graph to train ML models. |
| 44 | What is a 'Node Embedding' (such as FastRP)?                                 | A numerical representation of a node's local topology, often used as a feature in machine learning.                      |
| 45 | In an ML Link Prediction pipeline, what is the purpose of the 'Split' phase? | To divide the graph into disjoint sets for training, testing, and validation.                                            |
| 46 | What is the primary function of a 'Metadata Knowledge Graph'?                | To act as an enterprise-wide map recording the location, shape, and provenance of data assets.                           |
| 47 | In a metadata graph, what does a 'Data Sink' represent?                      | A terminal consumer of data, such as a dashboard or an ML feature set.                                                   |
| 48 | Term: Data Lineage                                                           | Definition: The tracking of data origins, movements, and transformations across various systems and tasks.               |
| 49 | What is 'Entity Resolution' (or Identity Resolution)?                        | The process of determining if two separate data records refer to the same real-world entity.                             |
| 50 | In identity graphs, why are 'Blocking Keys' used?                            | To reduce the search space by only comparing records that share a common high-level characteristic.                      |

| No | Question                                                                          | Answer                                                                                                                             |
|----|-----------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| 51 | What is the significance of a 'Golden Record' in master data management?          | It is the single, trusted, and deduplicated 'master' view of an entity created from multiple sources.                              |
| 52 | In the entity resolution process, what is a 'SAME_AS' relationship used for?      | To link nodes that have an unambiguous match on a strong identifier, like a Social Security number.                                |
| 53 | How does the 'Jaro-Winkler' distance assist in entity resolution?                 | It provides a string similarity metric used to compare weak identifiers like names.                                                |
| 54 | What is the purpose of 'EXPLAIN' in Cypher development?                           | It allows a developer to visualize a query plan to understand how the database will execute the query without actually running it. |
| 55 | How does 'PROFILE' differ from 'EXPLAIN' in Cypher?                               | PROFILE actually executes the query and provides real-time instrumentation data on performance and database hits.                  |
| 56 | In GDS production, why are 'Secondary' server instances used?                     | To isolate heavy computational workloads from the transactional processing occurring on 'Primary' servers.                         |
| 57 | What does a 'Causal Barrier' (or Bookmark) ensure in a distributed Neo4j cluster? | It guarantees that a user will always see at least their own previous writes, even when querying different servers.                |
| 58 | In a metadata graph, what is the 'Task' entity used to model?                     | Any discrete data job that processes or transforms a data asset.                                                                   |

| No | Question                                                                                             | Answer                                                                                                                |
|----|------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
| 59 | Which algorithm class would you use to predict if two users in a social network should be friends?   | Link Prediction.                                                                                                      |
| 60 | The property of 'ACID' transactions that ensures all parts of a transaction succeed or none do is _. | Atomicity                                                                                                             |
| 61 | Term: Raft                                                                                           | Definition: The consensus algorithm used by Neo4j to maintain synchronized transaction logs across a cluster.         |
| 62 | Why is property access in graphs considered linear time ( $O(N)$ )?                                  | Properties are stored as linked lists, requiring the database to scan the list to find a specific key.                |
| 63 | In the Cypher <code>UNWIND</code> clause, what is the result of applying it to a list?               | It transforms a single list of values into individual rows for subsequent query steps.                                |
| 64 | What is the 'Cold Start' problem in recommendation systems?                                          | The challenge of providing recommendations when there is little or no historical interaction data for a user or item. |
| 65 | What is the function of the <code>WHERE</code> clause in Cypher?                                     | It adds logical predicates to a <code>MATCH</code> pattern to filter results based on properties or relationships.    |
| 66 | A relationship that goes from one node back to the same node is called a _.                          | Self-relationship (or loop)                                                                                           |



| No | Question                                                                          | Answer                                                                                                                 |
|----|-----------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|
| 67 | In the context of the 'Data Importer', what is the role of an 'ID column'?        | It serves as a temporary unique identifier used to link nodes and relationships during the bulk import process.        |
| 68 | What does 'Index-Free Adjacency' eliminate the need for during a traversal?       | Global join indices or recursive lookups.                                                                              |
| 69 | Term: Pathfinding Algorithms                                                      | Definition: A class of algorithms used to find the most efficient route (shortest, fastest, cheapest) between nodes.   |
| 70 | In an Identity Knowledge Graph, what is 'Data Preparation'?                       | The phase where heterogeneous data is cleansed, standardized, and harmonized to make fields comparable.                |
| 71 | In Neo4j, what happens to a relationship if one of its endpoint nodes is deleted? | The database prevents the deletion unless the relationship is also explicitly deleted, as relationships cannot dangle. |
| 72 | How does 'Dijkstra's Algorithm' calculate the shortest path?                      | It explores the graph by evaluating the cumulative weights of relationship properties (like distance).                 |
| 73 | What is the primary benefit of using 'GraphQL' with a knowledge graph?            | It provides a strictly typed API schema that allows clients to request exactly the graph data they need.               |
| 74 | In GDS, what is the purpose of the 'Graph Catalog'?                               | It manages and stores named in-memory graph projections for reuse across different algorithmic tasks.                  |

| No | Question                                                                             | Answer                                                                                                                           |
|----|--------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| 75 | What is the risk of an 'Overly Ambitious Metamodel' early in a project?              | It leads to expensive, slow efforts that may become out-of-date before providing business value.                                 |
| 76 | Term: Hypergraph                                                                     | Definition: A graph model where a single relationship (edge) can connect more than two nodes.                                    |
| 77 | In Cypher, what is the difference between <code>SET</code> and <code>REMOVE</code> ? | <code>SET</code> adds or updates properties and labels, while <code>REMOVE</code> deletes them without deleting the node itself. |
| 78 | How can 'Weak Identifiers' be used to link entities safely?                          | By combining them into similarity scores and applying correcting factors based on other non-identifying features.                |
| 79 | Why is 'Context' considered the defining feature of a knowledge graph?               | Context is provided by the associative network of relationships and metadata that defines how entities interrelate.              |
| 80 | What occurs during the 'Execute' phase of the GDS workflow?                          | The chosen graph algorithm runs against the loaded in-memory projection to compute scores or patterns.                           |

## Flashcards hard

| No | Question                   | Answer                                                                                                                                                        |
|----|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1  | Knowledge Graph Definition | Interlinked sets of facts describing real-world entities and their interrelations in a human- and machine-understandable format with an organizing principle. |

| No | Question                                                                                            | Answer                                                                                                                                             |
|----|-----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|
| 2  | In the context of knowledge graphs, what is the primary role of an 'organizing principle'?          | It provides a layer of structure and context that enables reasoning and discovery over the underlying data.                                        |
| 3  | How does the 'organizing principle' function as a contract between a knowledge graph and its users? | It allows consuming software to expect specific structures, like labeled nodes and typed relationships, without needing detailed domain knowledge. |
| 4  | Term: Labeled Property Graph Model                                                                  | A graph model consisting of nodes, labels, relationships, and properties used to build high-fidelity domain representations.                       |
| 5  | In the property graph model, what primitive is used to represent an entity?                         | A node.                                                                                                                                            |
| 6  | Within a property graph, what do 'labels' signify on a node?                                        | The node's specific role or category within the graph domain.                                                                                      |
| 7  | What is the rule regarding 'dangling' relationships in a property graph model?                      | Relationships must always connect a start node and an end node.                                                                                    |
| 8  | In a property graph, where can 'properties' (key-value pairs) be stored?                            | Properties can be stored on both nodes and relationships.                                                                                          |
| 9  | How does 'index-free adjacency' allow for constant time ( $O(1)$ ) traversals?                      | It uses physical offsets (pointer-chasing) to move between nodes rather than global indexes.                                                       |

| No | Question                                                                                               | Answer                                                                                                                          |
|----|--------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| 10 | Term: Taxonomy                                                                                         | A classification scheme that organizes categories into a broader-narrower or generalization-specialization hierarchy.           |
| 11 | What is the primary structural difference between a taxonomy and an ontology?                          | Ontologies support complex, multi-level relationships and horizontal associativity beyond simple hierarchies.                   |
| 12 | The principle of 'just-enough semantics' advocates for what development approach?                      | Implementing semantic metadata incrementally based on current use cases rather than building an all-encompassing model upfront. |
| 13 | In graph database architecture, what is the 'causal barrier'?                                          | A mechanism ensuring that clients always see at least their own writes across a distributed cluster.                            |
| 14 | Which Cypher keyword is used to ensure a pattern exists without creating duplicates?                   | MERGE                                                                                                                           |
| 15 | What happens if a MERGE statement only partially matches an existing pattern in the graph?             | The database fails to match the pattern and creates a new instance of the entire pattern, potentially resulting in duplicates.  |
| 16 | To remove a node along with all its connected relationships in Cypher, use the <code>_</code> command. | DETACH DELETE                                                                                                                   |
| 17 | What is the primary design goal of the Cypher query language?                                          | To be a humane, expressive, and predominantly visual pattern-matching language.                                                 |

| No | Question                                                                                     | Answer                                                                                                                                            |
|----|----------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|
| 18 | In Cypher syntax, what do parentheses <code>()</code> represent?                             | Nodes                                                                                                                                             |
| 19 | How is a directed relationship with a type specified in Cypher ASCII art?                    | <code>-[:TYPE]-&gt;</code>                                                                                                                        |
| 20 | Concept: Graph Local Query                                                                   | A query anchored to specific known nodes, exploring the immediate or deep neighborhood of those nodes.                                            |
| 21 | What distinguishes a 'Graph Global Query' from a local one?                                  | Global queries process large parts or the entirety of the knowledge graph to generate aggregates without being bound to a specific starting node. |
| 22 | Cypher syntax:<br><code>[[:FRIEND*1..2]</code>                                               | Matches a relationship of type FRIEND at a path depth between one and two hops.                                                                   |
| 23 | Why is 'parameterization' essential when sending queries via a database driver?              | It allows the database to reuse query plans and protects the system against injection attacks.                                                    |
| 24 | What is the purpose of the Cypher <code>EXPLAIN</code> keyword?                              | It visualizes the query plan to help developers minimize 'DB hits' without executing the query.                                                   |
| 25 | How does <code>PROFILE</code> differ from <code>EXPLAIN</code> in Cypher performance tuning? | PROFILE executes the query and instruments it to show actual runtime behavior and execution costs.                                                |

| No | Question                                                                                               | Answer                                                                                                                     |
|----|--------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| 26 | In Neo4j internals, where are the 'neighbors' of a node essentially stored?                            | They are stored in a relationship chain directly referenced by the node, acting as a local index.                          |
| 27 | Why are property stores in Neo4j considered linear time ( $O(N)$ ) for reads?                          | Retrieving a specific property requires iterating through a list of property records attached to the node or relationship. |
| 28 | What role does the 'write-ahead log' play in Neo4j's ACID transactions?                                | It ensures durability by persisting intended updates to disk in an ordered fashion before applying them to the graph.      |
| 29 | The 'Raft' algorithm is utilized in Neo4j clusters to achieve what purpose?                            | To maintain transaction log consistency and synchronization across multiple servers.                                       |
| 30 | Which bulk-loading tool is best for high-speed, offline initial bootstrapping of billions of records?  | neo4j-admin import                                                                                                         |
| 31 | What is the function of the <code>UNWIND</code> clause in Cypher ingestion scripts?                    | It transforms a list of values into individual rows to be processed by subsequent statements.                              |
| 32 | In the <code>LOAD CSV</code> command, what is the purpose of the <code>WITH HEADERS</code> sub-clause? | It allows the query to reference data columns by their header names rather than by numerical index.                        |
| 33 | Term: Data Fabric                                                                                      | An organization-wide data access layer providing a connected, curated view of data across disparate underlying systems.    |

| No | Question                                                                                             | Answer                                                                                                                       |
|----|------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| 34 | What is a 'Composite Database' in the context of graph federation?                                   | A virtual database that provides a single point of access to multiple physical graph data sources.                           |
| 35 | How can 'virtual nodes' be created from external SQL or JSON sources without copying data?           | By using procedures from the APOC library to map external results into the Cypher query context.                             |
| 36 | What does the <code>@relationship</code> directive do in Neo4j's GraphQL implementation?             | It maps GraphQL fields to relationships between nodes in the underlying graph database.                                      |
| 37 | Category: Influence Algorithms (GDS)                                                                 | Algorithms that measure a node's importance or 'centrality' based on its position and connections in the network.            |
| 38 | What does 'Betweenness Centrality' calculate?                                                        | The number of shortest paths that pass through a specific node, identifying bridges and bottlenecks.                         |
| 39 | In Graph Data Science (GDS), what is a 'graph projection'?                                           | An in-memory, compressed representation of a specific subgraph optimized for high-performance algorithm execution.           |
| 40 | What is the operational difference between the GDS <code>mutate</code> and <code>write</code> modes? | Mutate saves results back into the in-memory projection, while Write persists them as properties in the underlying database. |

| No | Question                                                                                          | Answer                                                                                                                              |
|----|---------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| 41 | Concept: Community Detection                                                                      | Algorithms that partition a graph into sets of nodes that are more densely connected to each other than to the rest of the network. |
| 42 | In GDS, what is 'Topological Machine Learning'?                                                   | Using the graph's structure and connectivity patterns as the primary features for predictive modeling.                              |
| 43 | Term: Node Embedding                                                                              | A numerical vector that encodes a node's local topology and relative position within the graph.                                     |
| 44 | How does 'Graph Feature Engineering' improve machine learning models?                             | It extracts higher-quality features derived from graph topology, such as PageRank or community IDs, to train predictive models.     |
| 45 | In a GDS link prediction pipeline, what is the role of the 'combiner' (e.g., Cosine or Hadamard)? | It merges individual node properties into a single feature vector representing the potential relationship between a pair of nodes.  |
| 46 | Term: Metadata Knowledge Graph                                                                    | An enterprise map recording the location, schema, and provenance of data assets, processing tasks, and consumers.                   |
| 47 | What is a 'Data Sink' in a metadata graph model?                                                  | A terminal consumer of data, such as a dashboard or ML feature set, that does not produce a new dataset.                            |



| No | Question                                                                                              | Answer                                                                                                                                  |
|----|-------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| 48 | Term: Entity Resolution                                                                               | The process of determining whether two records from different data sources refer to the same real-world entity.                         |
| 49 | What is the purpose of 'Blocking Keys' in entity resolution?                                          | To reduce the search space by only comparing records that share a common attribute, preventing $n^2$ complexity.                        |
| 50 | In identity matching, how does a 'Strong Identifier' differ from a 'Weak' one?                        | Strong identifiers (like SSNs) provide unambiguous matches, while weak identifiers (like names) require probabilistic or fuzzy scoring. |
| 51 | What is the 'Jaro-Winkler' metric used for in data preparation?                                       | Measuring the similarity between two strings to identify potential duplicates in fuzzy matching rules.                                  |
| 52 | Concept: Golden Record                                                                                | A single, trusted master entity created by merging multiple resolved duplicate records from various source systems.                     |
| 53 | How does the 'FastRP' algorithm function in a GDS pipeline?                                           | It provides a fast method for generating node embeddings by projecting the graph's topology into a lower-dimensional vector space.      |
| 54 | In a metadata graph, what does 'Data Lineage' analysis track?                                         | The flow of data from its origin platforms through various transformation tasks to its final consumption points.                        |
| 55 | What logic does step 3 of the 'Graph-Based Entity Resolution' process apply to SIMILAR relationships? | It discards SIMILAR links if there is a known mismatch in strong identifier features like Passport Number or SSN.                       |

| No | Question                                                                                | Answer                                                                                                                                  |
|----|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| 56 | Why is relationship 'direction' mandatory when storing data using Cypher?               | It ensures a high-fidelity model by disambiguating the semantic relationship between nodes.                                             |
| 57 | What is the primary benefit of using 'Secondary Servers' in a Neo4j cluster for GDS?    | It provides physical isolation, preventing compute-intensive analytical workloads from impacting transactional performance.             |
| 58 | In GDS, what is the 'Alpha' tier for procedures?                                        | Experimental procedures that may change without notice and are not yet recommended for production use.                                  |
| 59 | How does 'Link Prediction' differ from 'Node Classification' in graph machine learning? | Link prediction forecasts missing connections between nodes, whereas node classification predicts the label or category of a node.      |
| 60 | Term: Semantic Search                                                                   | A search technique that uses an organizing principle to find relevant information based on meaning rather than literal string matching. |
| 61 | In GDS pipeline splitting, what is the 'Feature Set'?                                   | The portion of the graph reserved to compute the properties used as features for the machine learning model.                            |
| 62 | What is the function of the <code>gds.find_node_id</code> helper in the GDS Python API? | It retrieves the internal database ID for a node based on its labels and property values.                                               |

| No | Question                                                                                    | Answer                                                                                                                                |
|----|---------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| 63 | Within a metadata KG, what relationship typically connects a 'Dataset' to a 'DataPlatform'? | The source relationship.                                                                                                              |
| 64 | Term: Data Stewardship                                                                      | The management and oversight of an organization's data assets to ensure accessibility, quality, and compliance.                       |
| 65 | How can 'Betweenness Centrality' identify critical infrastructure in a transport network?   | By highlighting hubs where the failure of a single node would most significantly disrupt overall network connectivity.                |
| 66 | What is 'Causal Consistency' in a graph database?                                           | A consistency model that ensures the order of related operations is preserved across all nodes in a cluster.                          |
| 67 | In Cypher, what does the keyword <code>YIELD</code> signify when calling a procedure?       | It selects specific output columns from the procedure results to be used in the remainder of the query.                               |
| 68 | What is the role of a 'Driver' object in application-to-graph database connectivity?        | It manages the connection pool and protocol details for communicating with the database server.                                       |
| 69 | Cypher: <code>MATCH (n) DETACH DELETE n</code>                                              | Effectively deletes the entire graph by removing all nodes and their associated relationships.                                        |
| 70 | How does 'Graph Federation' via Neo4j Fabric handle cross-database queries?                 | It allows a single Cypher query to target multiple graphs by using the <code>USE</code> clause to specify different database aliases. |

| No | Question                                                                               | Answer                                                                                                                                               |
|----|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| 71 | Term: Semantic Bridge                                                                  | An ontology used to link and map concepts between different systems to a standard, well-understood vocabulary.                                       |
| 72 | In the context of GDS, what does 'Scale-Up' vs 'Scale-Out' imply for graph processing? | Scale-up uses a single large-memory machine for locality benefits, whereas scale-out distributes data across servers at the cost of network latency. |
| 73 | Why is 'Harmonization' a critical step in the Identity KG data preparation phase?      | It aligns disparate data formats (e.g., age vs. date of birth) into a common representation to make them comparable.                                 |
| 74 | What does a <code>SAME_AS</code> relationship indicate in an identity knowledge graph? | An unambiguous, high-confidence link between two records based on a shared strong identifier.                                                        |
| 75 | In GDS <code>train</code> results, what does 'AUCPR' measure?                          | The Area Under the Precision-Recall curve, used to evaluate the quality of a predictive model.                                                       |
| 76 | How does the 'Louvain' algorithm identify communities?                                 | It optimizes modularity by iteratively grouping nodes to maximize the density of internal connections versus external ones.                          |
| 77 | What is the primary risk of 'Boiling the Ocean' in knowledge graph modeling?           | Wasting time and resources capturing an overly ambitious, monolithic metamodel that may be out of date before completion.                            |

| No | Question                                                              | Answer                                                                                                                                      |
|----|-----------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| 78 | Cypher: <code>RETURN count{ (d)&lt;-[[:consumes]-(:DataSink) }</code> | An inline subquery used to count the number of DataSink nodes that have an incoming 'consumes' relationship from node <i>d</i> .            |
| 79 | In GDS, what is the purpose of the 'Graph Catalog'?                   | It manages named graph projections in memory so they can be reused across multiple algorithm executions.                                    |
| 80 | What is 'Index-Free Adjacency'?                                       | A storage technique where nodes directly store the memory addresses of their adjacent relationships, bypassing the need for global lookups. |

**Source:**